SURV703 – Content Analysis
"Testing" WEKA, explaining the results
Elisabeth Linek

Task: Install the appropriate version of WEKA (Links zu einer externen Webseite.) for your system. Convert the dataset into the ARFF format required by WEKA. Preprocess it with settings of your choice in the WEKA Explorer. Finally, perform a binary sentence classification (as shown in this week's lesson) with algorithms of your choice in the WEKA Experimenter. Hand in a table with the results and a brief interpretation in textual form.

First of all I followed the examples from the video lecture in order to have a "guided walk" through the WEKA GUI.

Starting with the preparation of the csv-file, in order get it into a WEKA-acceptable format, I set up the following:

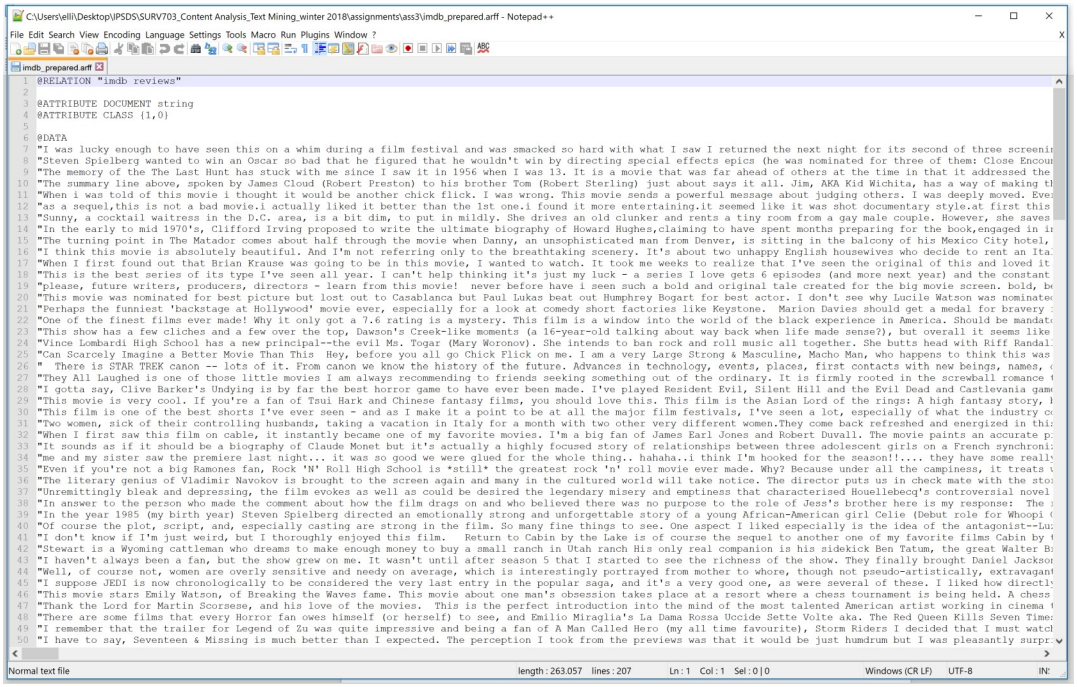

I saved the text-file with the extension *.arff (imdb_bow.arff)*, in order to import it as data file into the WEKA GUI. After loading the data file into WEKA, I followed the explained filter adoptions, changing the preselection into the following:

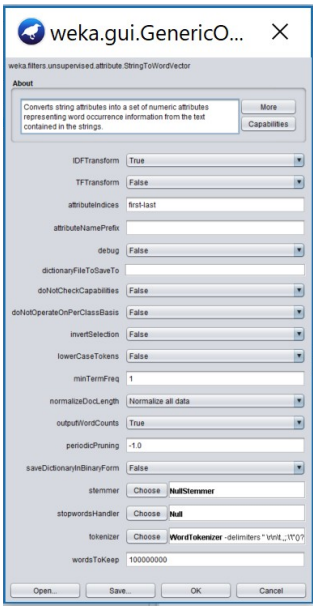

After the pre-selection was changed, the filter was applied to the DOCUMENT, and I went on to define or set the CLASS argument as attribute for the planned "implementation" of algorithms.
As final step of the preparation I saved the vectorized and normalized data set, giving it the name. *imdb_prepared_bow.arff*

———

Now starting with the application of the algorithms, I will jump to discuss briefly the results.

As shown within the video lectures I went on with a J48-tree algorithm. I changed the percentage of trainign data first to 75% and in a second step to 80%, just to find out what impact such a change might have on the results:

Scheme: weka.classifiers.trees.J48 -C 0.25 -M 2
[...]
Instances: 200
Attributes: 8153
Test mode: split 75.0% train, remainder test

J48 pruned tree

```
awful <= 0
|  worst <= 1.967162
|  |  entire <= 1.022614
|  |  |  nothing <= 1.83166
|  |  |  |  20 <= 1.056217
|  |  |  |  |  John <= 0
|  |  |  |  |  |  Some <= 0
|  |  |  |  |  |  |  family <= 1.554192
|  |  |  |  |  |  |  |  & <= 0
|  |  |  |  |  |  |  |  2 <= 0
|  |  |  |  |  |  |  |  |  away <= 2.129845
|  |  |  |  |  |  |  |  |  |  used <= 0
|  |  |  |  |  |  |  |  |  |  |  before <= 2.105742
|  |  |  |  |  |  |  |  |  |  |  |  minutes <= 1.992505
|  |  |  |  |  |  |  |  |  |  |  |  |  however <= 1.268507
|  |  |  |  |  |  |  |  |  |  |  |  |  |  and <= 0.060602
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  a <= 0.473168: 0 (5.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  a > 0.473168: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  and > 0.060602
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  were <= 2.622447
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  comedy <= 2.572386: 1 (68.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  comedy > 2.572386
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  It <= 1.069835: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  It > 1.069835: 0 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  were > 2.622447
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  This <= 0.718081: 0 (3.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  This > 0.718081: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  however > 1.268507
|  |  |  |  |  |  |  |  |  |  |  |  |  |  But <= 0: 0 (3.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  But > 0: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  minutes > 1.992505: 0 (5.0/1.0)
|  |  |  |  |  |  |  |  |  |  |  before > 2.105742: 0 (5.0/1.0)
|  |  |  |  |  |  |  |  |  |  used > 0
|  |  |  |  |  |  |  |  |  |  |  However <= 0: 0 (5.0)
|  |  |  |  |  |  |  |  |  |  |  However > 0: 1 (2.0)
|  |  |  |  |  |  |  |  |  away > 2.129845: 0 (6.0/1.0)
|  |  |  |  |  |  |  |  2 > 0: 0 (6.0/1.0)
|  |  |  |  |  |  |  & > 0: 1 (7.0)
|  |  |  |  |  |  family > 1.554192: 0 (8.0/1.0)
|  |  |  |  |  Some > 0: 0 (7.0/1.0)
|  |  |  |  John > 0: 1 (8.0)
|  |  |  20 > 1.056217: 0 (8.0/1.0)
|  |  nothing > 1.83166: 0 (10.0)
|  entire > 1.022614: 0 (9.0)
|  worst > 1.967162: 0 (11.0)
awful > 0: 0 (14.0)
```

Number of Leaves: 24
Size of the tree: 47

=== Summary ===

Correctly Classified Instances	27	54	%
Incorrectly Classified Instances	23	46	%
Kappa statistic	0.077		
Mean absolute error	0.4573		
Root mean squared error	0.6606		
Relative absolute error	91.2745 %		
Root relative squared error	131.7917 %		
Total Number of Instances	50		

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0,522	0,444	0,500	0,522	0,511	0,077	0,553	0,491	1
	0,556	0,478	0,577	0,556	0,566	0,077	0,553	0,575	0
Weighted Avg.	0,540	0,463	0,542	0,540	0,541	0,077	0,553	0,536	

=== Confusion Matrix ===

```
 a  b   <-- classified as
12 11 |  a = 1
12 15 |  b = 0
```

Followed by a J48 tree algorithm based on 80% of trainign data from the data set:

=== Run information ===

Scheme: weka.classifiers.trees.J48 -C 0.25 -M 2
[...]
Instances: 200
Attributes: 8153
Test mode: split 80.0% train, remainder test

J48 pruned tree

```
------------------
awful <= 0
|  worst <= 1.967162
|  |  entire <= 1.022614
|  |  |  nothing <= 1.83166
|  |  |  |  20 <= 1.056217
|  |  |  |  |  John <= 0
|  |  |  |  |  |  Some <= 0
|  |  |  |  |  |  |  family <= 1.554192
|  |  |  |  |  |  |  |  & <= 0
|  |  |  |  |  |  |  |  |  2 <= 0
|  |  |  |  |  |  |  |  |  |  away <= 2.129845
|  |  |  |  |  |  |  |  |  |  |  used <= 0
|  |  |  |  |  |  |  |  |  |  |  |  before <= 2.105742
|  |  |  |  |  |  |  |  |  |  |  |  |  minutes <= 1.992505
|  |  |  |  |  |  |  |  |  |  |  |  |  |  however <= 1.268507
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  and <= 0.060602
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  a <= 0.473168: 0 (5.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  a > 0.473168: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  and > 0.060602
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  were <= 2.622447
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  comedy <= 2.572386: 1 (68.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  comedy > 2.572386
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  It <= 1.069835: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  It > 1.069835: 0 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  were > 2.622447
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  This <= 0.718081: 0 (3.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  This > 0.718081: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  however > 1.268507
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  But <= 0: 0 (3.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  But > 0: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  |  minutes > 1.992505: 0 (5.0/1.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  before > 2.105742: 0 (5.0/1.0)
|  |  |  |  |  |  |  |  |  |  |  |  used > 0
|  |  |  |  |  |  |  |  |  |  |  |  |  However <= 0: 0 (5.0)
|  |  |  |  |  |  |  |  |  |  |  |  |  However > 0: 1 (2.0)
|  |  |  |  |  |  |  |  |  |  |  away > 2.129845: 0 (6.0/1.0)
|  |  |  |  |  |  |  |  |  |  2 > 0: 0 (6.0/1.0)
|  |  |  |  |  |  |  |  |  & > 0: 1 (7.0)
|  |  |  |  |  |  |  |  family > 1.554192: 0 (8.0/1.0)
|  |  |  |  |  |  |  Some > 0: 0 (7.0/1.0)
|  |  |  |  |  |  John > 0: 1 (8.0)
|  |  |  |  |  20 > 1.056217: 0 (8.0/1.0)
|  |  |  |  nothing > 1.83166: 0 (10.0)
|  |  |  entire > 1.022614: 0 (9.0)
|  |  worst > 1.967162: 0 (11.0)
awful > 0: 0 (14.0)
```

Number of Leaves: 24
Size of the tree: 47

=== Summary ===

Correctly Classified Instances	20	50 %
Incorrectly Classified Instances	20	50 %
Kappa statistic	0.0123	
Mean absolute error	0.5074	
Root mean squared error	0.6868	
Relative absolute error	101.4189 %	
Root relative squared error	137.2681 %	
Total Number of Instances	40	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0,632	0,619	0,480	0,632	0,545	0,013	0,514	0,499	1
	0,381	0,368	0,533	0,381	0,444	0,013	0,514	0,521	0
Weighted Avg.	0,500	0,487	0,508	0,500	0,492	0,013	0,514	0,511	

=== Confusion Matrix ===

```
 a  b   <-- classified as
12  7 |  a = 1
13  8 |  b = 0
```

→ I have to admit, that I was surprised by the difference, just based on the change of the amount of training data compared to the remaining test data. The correctly classified instances differ from 54% to 50%. The confusion matrix differs as well.

What surprised me most is the fact, that the level of precision differs much from the example we saw on the video lectures, where we reached levels around 90%. That will depend on the provided data set, but I would like to know which other filters I could have set to reach better results.

As another test I went on with a Naive Bayes algorithm:

=== Run information ===

Scheme: weka.classifiers.bayes.NaiveBayes
[...]
Instances: 200
Attributes: 8153
Test mode: split 80.0% train, remainder test

Naive Bayes Classifier (just some examples, the whole result overview would have needed much more space)

Attribute	Class 1 (0.5)	Class 0 (0.5)
$1		
mean	0.0187	0
std. dev.	0.3113	0.3113
weight sum	100	100
precision	1.868	1.868
&		
mean	0.6253	0.0853
std. dev.	2.3977	0.6298
weight sum	100	100
precision	1.4212	1.4212
*		
mean	0.0374	0
std. dev.	0.6227	0.6227
weight sum	100	100
precision	3.7361	3.7361
Jack		
mean	0.0423	0.1479
std. dev.	0.3521	1.0877
weight sum	100	100
precision	2.1125	2.1125
friend		
mean	0.1225	0.2205
std. dev.	0.5084	1.0948
weight sum	100	100
precision	0.8167	0.8167
friends		
mean	0.217	0.1346
std. dev.	0.8292	0.6458
weight sum	100	100
precision	0.434	0.434
frightening		
mean	0.0803	0.0357
std. dev.	0.4567	0.3552
weight sum	100	100
precision	0.8924	0.8924
yes		
mean	0	0.0887
std. dev.	0.2956	0.5814
weight sum	100	100
precision	1.7734	1.7734
zombie/cannibal		
mean	0	0.0764
std. dev.	1.2741	1.2741
weight sum	100	100
precision	7.6448	7.6448

=== Summary ===

Correctly Classified Instances	22	55 %
Incorrectly Classified Instances	18	45 %
Kappa statistic	0.0932	
Mean absolute error	0.4537	
Root mean squared error	0.6712	
Relative absolute error	90.6865 %	
Root relative squared error	134.1532 %	
Total Number of Instances	40	

=== Detailed Accuracy By Class ===

TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0,474	0,381	0,529	0,474	0,500	0,094	0,590	0,539	1
0,619	0,526	0,565	0,619	0,591	0,094	0,564	0,552	0

Weighted Avg.	0,550	0,457	0,548	0,550	0,548	0,094	0,576	0,546

=== Confusion Matrix ===

```
 a  b   <-- classified as
 9 10 |  a = 1
 8 13 |  b = 0
```

The results highlighted here are close to the results I reached with the J48 algorithm, based on 75% training data. The confusion matrix showing, just as the correction level already did, that the algorithm are not leading to a very successful automatic classification. I am still wondering what I depends on. Maybe the normalisation of the original data was not fitting well?

Leaving the classification based WEKA GUI behind, turning to the experimental tool, in order to compare the algorithms, I got the following results.
In a first step I selected the following algorithms which lead to an overload of m system, and I interrupted the test.

- ZeroR: No model built yet.
- Naive Bayes Classifier: No model built yet.
- SMO: No model built yet.
- IBk: No model built yet.
- J48
- Decision Stump: No model built yet. (left aside in the second trial)

I decided to work according to the example from lecture: 10 repetitions each based on 10% of the data.

The results:
18:03:21: Started
20:14:03: User aborting experiment.
20:14:51: Interrupted
20:14:51: There were 0 errors

I went on with a second try, downgrading the number of folds to 5 (even though it is not the prefered method), the second trial was successful.

1) Percent of correctness selected: (according to video lectures)

```
Tester:     weka.experiment.PairedCorrectedTTester -G 4,5,6 -D 1 -R 2 -S 0.05 -result-matrix [...]
Analysing:  Percent_correct
Datasets:   1
Resultsets: 5
Confidence: 0.05 (two tailed)
Sorted by:  -
Date:       13.12.18 20:32

Dataset                    (1) rules.Ze | (2) bayes (3) funct (4) lazy.   (5) trees
--------------------------------------------------------------------------------
'imdb reviews-weka.filter (50)   50.00 |   65.25 v       74.55 v   59.35 v   62.05 v
--------------------------------------------------------------------------------
                                 (v/ /*) |   (1/0/0)   (1/0/0)   (1/0/0)   (1/0/0)


Key:
(1) rules.ZeroR " 48055541465867954
(2) bayes.NaiveBayes " 5995231201785697655
(3) functions.SMO '-C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K \"functions.supportVector.PolyKernel -E 1.0 -C 250007\" -
calibrator \"functions.Logistic -R 1.0E-8 -M -1 -num-decimal-places 4\"' -6585883636378691736
(4) lazy.IBk '-K 1 -W 0 -A \"weka.core.neighboursearch.LinearNNSearch -A \\\"weka.core.EuclideanDistance -R first-last\\\"\"' -
3080186098777067172
(5) trees.J48 '-C 0.25 -M 2' -217733168393644444
```

→ The comparison of the given levels of correct classifications show, that the “function SMO” algorithm reached the highest level.

2) Precision selected: (as shown in video)

Tester: weka.experiment.PairedCorrectedTTester -G 4,5,6 -D 1 -R 2 -S 0.05 -result-matrix "weka.experiment.ResultMatrixPlainText -mean-prec 2 -stddev-prec 2 -col-name-width 0 -row-name-width 25 -mean-width 2 -stddev-width 2 -sig-width 1 -count-width 5 -print-col-names -print-row-names -enum-col-names"
Analysing: IR_precision
Datasets: 1
Resultsets: 5
Confidence: 0.05 (two tailed)
Sorted by: -
Date: 13.12.18 20:35

Dataset	(1) rules.Z	(2) baye	(3) func	(4) lazy	(5) tree
'imdb reviews-weka.filter (50)	0.50	0.68 v	0.73 v	0.58 v	0.63 v
(v/ /*)		(1/0/0)	(1/0/0)	(1/0/0)	(1/0/0)

Key:
(1) rules.ZeroR " 48055541465867954
(2) bayes.NaiveBayes " 5995231201785697655
(3) functions.SMO '-C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K \"functions.supportVector.PolyKernel -E 1.0 -C 250007\" -calibrator \"functions.Logistic -R 1.0E-8 -M -1 -num-decimal-places 4\"' -6585883636378691736
(4) lazy.IBk '-K 1 -W 0 -A \"weka.core.neighboursearch.LinearNNSearch -A \\\"weka.core.EuclideanDistance -R first-last\\\"\"' -3080186098777067172
(5) trees.J48 '-C 0.25 -M 2' -217733168393644444

→ The comparison of precision levels shows as well, that the “function SMO” algorithm reaches the highest level of correct classifications. But still, we do not reach a level of 80% or higher, as we saw in the video lectures.

3) F measure selected: (as shown in video lectures)

Tester: weka.experiment.PairedCorrectedTTester -G 4,5,6 -D 1 -R 2 -S 0.05 -result-matrix "weka.experiment.ResultMatrixPlainText -mean-prec 2 -stddev-prec 2 -col-name-width 0 -row-name-width 25 -mean-width 2 -stddev-width 2 -sig-width 1 -count-width 5 -print-col-names -print-row-names -enum-col-names"
Analysing: F_measure
Datasets: 1
Resultsets: 5
Confidence: 0.05 (two tailed)
Sorted by: -
Date: 13.12.18 20:36

Dataset	(1) rules.Z	(2) baye	(3) func	(4) lazy	(5) tree
'imdb reviews-weka.filter (50)	0.67	0.62	0.75 v	0.65	0.61
(v/ /*)		(0/1/0)	(1/0/0)	(0/1/0)	(0/1/0)

Key:
(1) rules.ZeroR " 48055541465867954
(2) bayes.NaiveBayes " 5995231201785697655
(3) functions.SMO '-C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K \"functions.supportVector.PolyKernel -E 1.0 -C 250007\" -calibrator \"functions.Logistic -R 1.0E-8 -M -1 -num-decimal-places 4\"' -6585883636378691736
(4) lazy.IBk '-K 1 -W 0 -A \"weka.core.neighboursearch.LinearNNSearch -A \\\"weka.core.EuclideanDistance -R first-last\\\"\"' -3080186098777067172
(5) trees.J48 '-C 0.25 -M 2' -217733168393644444

Overview of analysis:

- 20:32:13 - Available resultsets
- 20:32:48 - Percent_correct - rules.ZeroR " 48055541465867954
- 20:35:21 - IR_precision - rules.ZeroR " 48055541465867954
- 20:36:44 - F_measure - rules.ZeroR " 48055541465867954

→ The comparison of the f test levels shows what we could see before, the "function SMO" algorithm reaches the highest level regarding the classifications.
But still, we do not reach the level we saw in the video lectures.

I would assume, that further data preparation at the beginning of the planned classification could lead to higher or better results, with less confusion within the classifications.